

Investment Memorandum

Sarvam AI Technologies Pvt. Ltd. — Series B Growth Round

Prepared by:	Deal Team, Together Fund	Round:	Series B (growth)
Date:	July 9, 2026	Post-money valuation:	\$1.5B (reported)
Sector:	Applied AI / Foundational Models — India	Round size:	\$300M target (\$234M first close)
Recommendation:	Proceed to partner vote — pursue allocation	Stage:	Growth / Series B

1. Executive Summary

Sarvam AI is a Bengaluru-based full-stack artificial intelligence company building sovereign, multilingual foundation models and speech technology purpose-built for India's 22 scheduled languages. Founded in mid-2023 by **Vivek Raghavan** (co-creator of India's Aadhaar biometric identity system at UIDAI) and **Pratyush Kumar** (co-founder of the AI4Bharat research group at IIT Madras), the company has moved from a \$41M combined seed and Series A in December 2023 to a reported \$1.5B post-money Series B in June 2026 — a roughly 37x step-up in under three years, one of the fastest valuation trajectories of any Indian AI company to date.

The Series B, targeting \$300M, had a first close of \$234M as of June 15, 2026, led by **HCLTech** (committing \$150M for a reported ~10.5% stake, its largest strategic AI investment to date) with continued participation from **Bessemer Venture Partners**, **Khosla Ventures**, and **Peak XV Partners** (formerly Sequoia India). The remaining book is reportedly still being allocated, which is the basis for this memo: an evaluation of whether Together Fund should seek an allocation in the remaining ~\$66M of the round.

Sarvam's most distinguishing asset is not a single model but a government relationship: in April 2025 it was selected by India's Ministry of Electronics and IT (MeitY) under the **IndiaAI Mission** to build an indigenous foundation model, with access to subsidized government-procured GPU compute — a structural cost advantage few private AI labs anywhere have. It has since shipped two model families (Sarvam-30B, a mixture-of-experts model, and Sarvam-105B / "Indus," a 105B-parameter, ~9B-active-parameter MoE model with a 128K context window) and signed enterprise deployments with State Bank of India (SBI) Life Insurance, LIC, IDFC First Bank, Tata Capital, and CRED.

We view Sarvam as a credible, if richly priced, way to gain exposure to the "sovereign AI" thesis — the idea that large, non-US markets with linguistically diverse populations and data residency concerns will not simply consume US frontier models, but will build (or co-build with government support) their own. The central open question is whether a \$1.5B valuation for a company that does not disclose revenue, and whose flagship customers are concentrated in Indian BFSI and government, is underwritten by anything more than narrative and compute access. Our recommendation, on balance, is to proceed to a partner vote in favor of a modest allocation, sized to reflect high conviction on the thesis and lower conviction on near-term monetization.

2. Company Overview

Sarvam AI Technologies Pvt. Ltd. was founded in mid-2023 (the company was described as "five months old" at the time of its December 2023 funding announcement, implying incorporation around July 2023) and is headquartered in

Bengaluru, Karnataka. The company describes its mission as building “the full stack of AI for India” — foundation models, speech and voice infrastructure, and applied enterprise products — rather than a single point product, on the premise that India’s language diversity (22 scheduled languages, hundreds of dialects, and a majority of the population that is far more comfortable in speech than in English text) is not adequately served by general-purpose Western LLMs.

Headcount has grown quickly alongside its funding: reported figures vary by source and point in time — approximately 114 employees in mid-2025, roughly 200 shortly after, and figures as high as 354 in the most recent press coverage tied to the Series B. We treat the exact current number as directionally rather than precisely known (a common issue with fast-growing private companies) but the trend — roughly 3x growth in about a year — is a reasonably reliable signal of scaling intent tied to the new capital and the IndiaAI Mission compute allocation.

3. Founders & Leadership Team

Vivek Raghavan — Co-founder & CEO

Best known as a co-creator and former Deputy Director General / Chief Product Manager of the biometric matching system underlying **Aadhaar**, India’s national digital identity program covering over a billion residents — one of the largest biometric deduplication systems ever built. That background is directly relevant to Sarvam’s pitch: it is difficult to name another founder anywhere who has shipped identity/AI infrastructure at the scale of the Indian population, and it gives Sarvam unusually credible access to government procurement processes, which is a meaningful part of the IndiaAI Mission selection and the enterprise BFSI customer base.

Pratyush Kumar — Co-founder

Co-founder of **AI4Bharat**, a language-AI research group housed at IIT Madras focused on NLP and speech technology for Indian languages, which is the direct research lineage for Sarvam’s multilingual model and speech stack. Prior to Sarvam, Kumar’s research background is in large-scale ML systems; the AI4Bharat connection also gives Sarvam a recruiting pipeline into India’s small pool of language-AI research talent, which is a genuine bottleneck for competitors without that lineage.

Diligence note: we have not independently verified a full executive bench (CTO/Head of Research/CFO) beyond the two co-founders from public sources reviewed for this memo; this should be confirmed directly in management diligence calls before final terms.

4. Product & Technology

Sarvam’s product surface spans three layers: foundation models, speech/voice infrastructure, and applied enterprise deployments built on top of both.

- **Sarvam-30B** — a mixture-of-experts (MoE) language model tuned for Indian-language understanding and generation, released as part of the company’s February 2026 model wave.
- **Sarvam-105B (“Indus”)** — the flagship model, ~105B total parameters with roughly 9B active parameters per token (MoE routing), a 128K token context window, and a beta release on February 20, 2026. Indus is positioned as India’s assistant-grade model for complex reasoning and agentic workflows, the closest public analogue to a GPT-4-class “national” model India has shipped domestically.

- **Multilingual coverage** — the company states support across all 22 languages scheduled in the Indian constitution, spanning both text and speech (ASR/TTS), which is the core technical moat versus general-purpose global models trained overwhelmingly on English and a handful of high-resource languages.
- **Enterprise deployment layer** — vertical-specific fine-tunes and integrations for banking, insurance, telecom, healthcare, and government use cases, delivered via API and on-premise/government-cloud deployment where data residency requires it.

Roadmap items called out alongside the Series B include a next frontier model targeting agentic AI, coding, and cybersecurity use cases — a signal that Sarvam intends to compete more directly with general-purpose frontier labs rather than staying purely in the multilingual niche, which raises the compute-intensity (and therefore burn) profile of the roadmap materially.

5. Government Partnership — the IndiaAI Mission

In April 2025, India's Ministry of Electronics and Information Technology (MeitY) selected Sarvam AI, under the government's **IndiaAI Mission**, to build an indigenous foundational model, with access to GPU compute procured and subsidized by the government. This is arguably the single most important asset in the company relative to peers: frontier-model training and inference are compute-bound businesses, and preferential access to subsidized compute directly addresses the largest cost line for any lab at this stage.

It is also, however, a source of key-person and key-relationship risk: government mission selections can be politically contingent, subject to re-tendering, or reallocated across successive administrations or budget cycles in ways that a purely commercial customer relationship would not be. We treat this as a genuine strength today and a monitorable risk over a multi-year hold.

6. Customers & Traction

Disclosed enterprise customers, drawn from public reporting rather than a data room, span several of India's largest financial institutions:

- **SBI Life Insurance** — insurance arm of State Bank of India, India's largest bank.
- **LIC (Life Insurance Corporation of India)** — India's largest life insurer, state-owned.
- **IDFC First Bank** — private-sector Indian bank.
- **Tata Capital** — financial services arm of the Tata Group.
- **CRED** — fintech / credit-card bill payment and rewards platform, one of India's highest-profile consumer fintech brands.

Sarvam does not publicly disclose revenue, ARR, or customer count beyond these named logos, which is typical for a private company at this stage but means this memo cannot independently verify unit economics, net revenue retention, or customer concentration risk — all of which should be primary asks in a management data-room request ahead of any final commitment. What can be said with more confidence is that the customer set is concentrated in regulated BFSI, which is consistent with a company whose core differentiation (data residency, government trust, language coverage) is most valuable exactly in those regulated, India-domestic verticals.

7. Market Opportunity

The bet underlying Sarvam is that India will not be a pure-consumption market for US frontier labs. Several structural factors support this: (1) India has 22 constitutionally scheduled languages and hundreds of millions of users who are far more comfortable in speech and vernacular text than in English, a segment global labs serve poorly; (2) the Indian government has explicit sovereign-AI and data-residency ambitions, expressed concretely through the IndiaAI Mission's subsidized-compute program; and (3) regulated sectors — banking, insurance, government services, defense — have genuine data-residency and auditability constraints that favor a domestic vendor with a government-endorsed track record over a foreign hyperscaler API.

Against this, the ceiling on a purely India-domestic AI business is a real constraint: global frontier labs (OpenAI, Google, Anthropic, Meta) are all actively localizing for Indian languages and pricing aggressively to win the same enterprise and government logos, and the addressable market for a \$1.5B-valued company to “grow into” is smaller than a comparable US or China-focused peer would enjoy at the same valuation, at least until/unless Sarvam's roadmap into agentic AI, coding, and cybersecurity succeeds in expanding its addressable market beyond language localization.

8. Competitive Landscape

Sarvam sits between two categories of competitor, neither of which is a clean comparable:

- **Global frontier labs** (OpenAI, Google, Anthropic, Meta) — vastly better resourced on compute and general model quality, increasingly investing in Indian-language support and local go-to-market, but structurally weaker on data residency guarantees and government trust.
- **Other Indian sovereign-AI initiatives** — including Reliance/Ola's Krutrim and several smaller BharatGPT-style efforts — pursuing an overlapping thesis with different capital structures (Krutrim in particular is backed by a large domestic conglomerate rather than a venture syndicate). Sarvam's differentiation versus these peers rests on its founders' government-scale delivery track record (Aadhaar) and its formal IndiaAI Mission selection, rather than on model benchmarks alone, which are difficult to verify independently and change quickly.

We were not able to independently verify head-to-head benchmark performance of Sarvam-105B against either global frontier models or domestic peers from public sources reviewed for this memo; any claims of model quality superiority should be treated as company-sourced until verified in technical diligence.

9. Funding History

Date	Round	Amount	Lead / Key Participants	Reported Valuation
Dec 2023	Seed + Series A\$41M (combined)		Series A led by Lightspeed; seed co-led by Peak XV Partners; Khosla Ventures participated	Not disclosed
Jun 15, 2026	Series B (first close)	\$234M of \$300M target	Led by HCLTech (\$150M); Bessemer Venture Partners; continued participation from Khosla Ventures and Peak XV Partners	\$1.5B post-money

Sources: company and investor disclosures reported in TechCrunch, Bloomberg, and Business Standard coverage of the respective rounds. The Series B was, at time of writing, still accepting allocation toward its \$300M target beyond the \$234M first close — the basis for Together Fund's evaluation of a follow-on or new-money position in this memo.

The step-up from an undisclosed seed/Series A valuation to a \$1.5B post-money in roughly two and a half years, and the involvement of a large strategic corporate investor (HCLTech) writing the largest single check in the round, both point to a deal being priced substantially on strategic and narrative grounds — government relationship, founder pedigree, and category positioning — rather than on disclosed financial metrics. This is not unusual for AI-infrastructure investing at this point in the cycle, but it does mean Together Fund's underwriting here is closer to a thesis bet than a multiple-on-comparable-revenue analysis.

10. Deal Terms Under Consideration

Terms below reflect the reported structure of the round as a going concern; final terms for any Together Fund allocation would need to be confirmed directly with the company or lead investors, as we do not have access to the term sheet at the time of this memo.

Item	Reported / Assumed Terms
Round	Series B, target \$300M, \$234M first close closed
Post-money valuation	\$1.5B (reported)
Lead investor	HCLTech — \$150M, reported ~10.5% resulting stake
Other participants	Bessemer Venture Partners, Khosla Ventures, Peak XV Partners
Remaining allocation	Approximately \$66M to reach the stated \$300M target
Proposed Together Fund ticket	\$3–5M, subject to partner vote and allocation availability
Instrument	To be confirmed — likely priced equity given strategic lead's size

11. Thesis Fit for Together Fund

Sarvam maps cleanly onto several planks of Together Fund's applied-AI and AI-infrastructure thesis: it is a foundation-model company building genuinely differentiated, non-commodity infrastructure (multilingual sovereign models) rather than a thin wrapper on an existing frontier API; it has real enterprise deployment traction in regulated verticals rather than purely consumer or developer-tool usage; and it has a structural, government-endorsed compute and procurement advantage that is very difficult for a new entrant to replicate. It is also a relevant proof point for the broader thesis that large non-US markets will develop indigenous AI champions rather than fully importing US models — a pattern we are tracking across several other geographies.

Where it fits less cleanly: at a \$1.5B valuation and growth stage, this is a later, larger check than Together Fund's typical entry point, and the round is being led by a strategic corporate investor rather than a financial sponsor, which can affect governance dynamics and future-round dynamics in ways worth understanding before committing capital.

12. Key Risks & Mitigants

Risk	Mitigant / Notes
Revenue and unit economics are not disclosed; valuation may be running ahead of monetization.	Request revenue, ARR, and customer-concentration data directly in diligence before final commit.
Government-mission dependency (IndiaAI Mission compute access) could be politically or budget-cycle contingent.	Assess contractual durability of the compute allocation and whether it survives a change in administration or ministry leadership.
Key-person concentration on two founders; full executive bench not independently verified.	Confirm CTO/Head of Research/CFO identities and tenure directly with the company.
Global frontier labs (OpenAI, Google, Anthropic, Meta) are actively localizing for Indian languages and could compress Sarvam's moat.	Track benchmark parity and pricing over the next 2–3 model release cycles.
Roadmap expansion into agentic AI, coding, and cybersecurity increases compute intensity and burn versus the original multilingual-model thesis.	Request updated burn multiple and runway projections reflecting the expanded roadmap.
Customer base concentrated in Indian BFSI and government; limited disclosed geographic or vertical diversification.	Evaluate whether pipeline outside BFSI/government (e.g., healthcare, telecom, agriculture, defense) is converting to signed contracts.

13. Open Diligence Questions

- What is current ARR (or equivalent revenue metric), and how much of it is government / IndiaAI Mission-linked versus purely commercial enterprise contracts?
 - What is the exact scope, duration, and renewal mechanism of the IndiaAI Mission compute allocation, and what happens to it in a change-of-government or re-tender scenario?
 - What is the full executive team beyond the two named co-founders, and what is founder equity and vesting status post-Series B?
 - How does Sarvam-105B / Indus perform on independent, third-party multilingual and reasoning benchmarks versus both global frontier models and domestic peers (e.g., Krutrim)?
 - What is customer concentration — do the five named BFSI logos represent the majority of revenue, and what is the sales pipeline outside BFSI and government?
 - What is the instrument, liquidation preference, and governance structure being offered to non-lead participants joining the remaining Series B allocation?
-

14. Recommendation

We recommend Together Fund proceed to a partner vote in favor of pursuing a modest allocation (indicatively \$3–5M) in the remaining tranche of Sarvam AI's Series B, subject to: (1) direct management diligence calls covering the open questions in Section 13, particularly revenue and the durability of the IndiaAI Mission relationship; (2) confirmation of instrument and governance terms available to non-lead participants; and (3) sizing the check to reflect that this is principally a thesis and category bet on sovereign, multilingual AI infrastructure in a government-endorsed champion, rather than a valuation underwritten by disclosed financial performance. Given the strength of the founder pedigree, the structural compute advantage, and the real (if narrow) enterprise traction, we believe the risk/reward supports participation at a size the fund can afford to be wrong about.

This memorandum was prepared for internal Together Fund deal-team use based on information reported publicly by the company, its investors, and financial media as of July 2026, together with reasonable analyst interpretation. It does not reflect data-room materials, management representations, or independent financial diligence, and should not be relied upon as a complete basis for an investment decision absent further diligence as outlined in Section 13.